

# HW4

## Bitcoin Miner

Introduction to Parallel Computing  
2022/04/26

# Spec

<https://hackmd.io/@ipc22/hw4>

# Introduction - What is a Miner?

- A miner's job is to find out a hash value, such that the block is secured
- More specifically, the primary task of a miner is to find out a proper value called “nonce”, which is part of a block's header
- The value of nonce has to satisfy some requirements, including:
  - It has to be 32-bit long
  - It has to be less than a predefined value called “Target Difficulty”

# Introduction - What is a Block?

- A block consists of two parts: the block header and block body
- The block body contains a number of bitcoin transactions (or simply transactions) and coinbase transactions
  - Bitcoin transactions are the records of bitcoin transfers
  - Coinbase transactions records the bitcoins rewarded to the miner who finds the proper nonce for the block
- There are several fields included in the block header. We only cares about the following fields.
  - Nonce
  - The hash value of the previous block
  - Merkle root
  - Target difficulty (nBits)

# Introduction - Block Layout

- The layout of a bitcoin block looks like the table presented below.
  - The yellow part corresponds to the block header
  - The gray part corresponds to the block body
- A secure hash value is generated (on the right side) for the block header

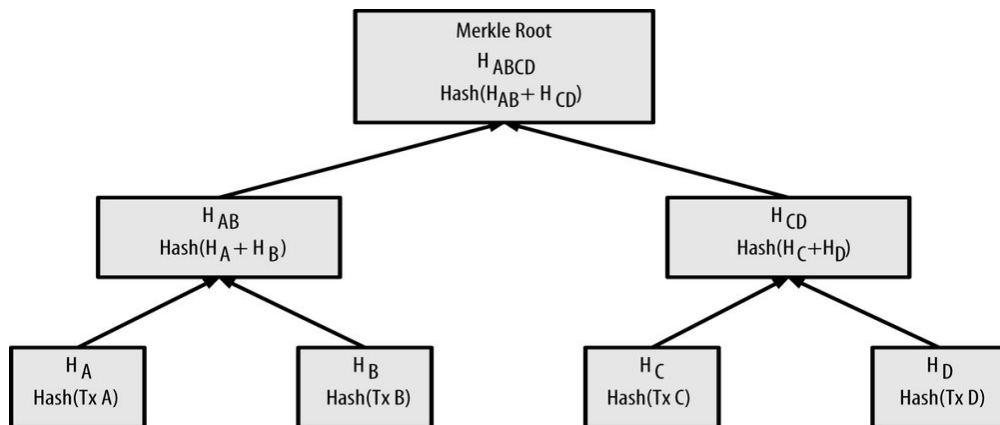
version	02000000
previous block hash (reversed)	17975b97c18ed1f7e255adf297599b55 330edab87803c8170100000000000000
Merkle root (reversed)	8a97295a2747b4f1a0b3948df3990344 c0e19fa6b2b92b3a19c8e6badc141787
timestamp	358b0553
bits	535f0119
nonce	48750833
transaction count	63
coinbase transaction	
transaction	
...	

Block hash

0000000000000000  
e067a478024addfe  
cdc93628978aa52d  
91fabd4292982a50

# Introduction - Merkle Tree Nodes

- We store the transactions in a data structure called “Merkle Tree”
  - The tree is constructed in a binary tree fashion
  - The root and intermediate nodes of the tree are generated from their children



$$\begin{aligned} H_A &= \text{SHA256}(\text{SHA256}(\text{TxA})) \\ H_B &= \text{SHA256}(\text{SHA256}(\text{TxB})) \\ H_{AB} &= \text{SHA256}(\text{SHA256}(H_A + H_B)) \end{aligned}$$

# Introduction - Double SHA-256

- Each parent node is generated by a double SHA-256 function
  - Each child node is encrypted by SHA-256 TWICE
  - The encrypted results (H<sub>A</sub> and H<sub>B</sub>) are summed together
  - The sum is then encrypted AGAIN by SHA-256 TWICE
- The procedure is repeated again and again, until the root node is reached
  - The root node is called the Merkle Root

$$H_A = \text{SHA256}(\text{SHA256}(\text{TxA}))$$



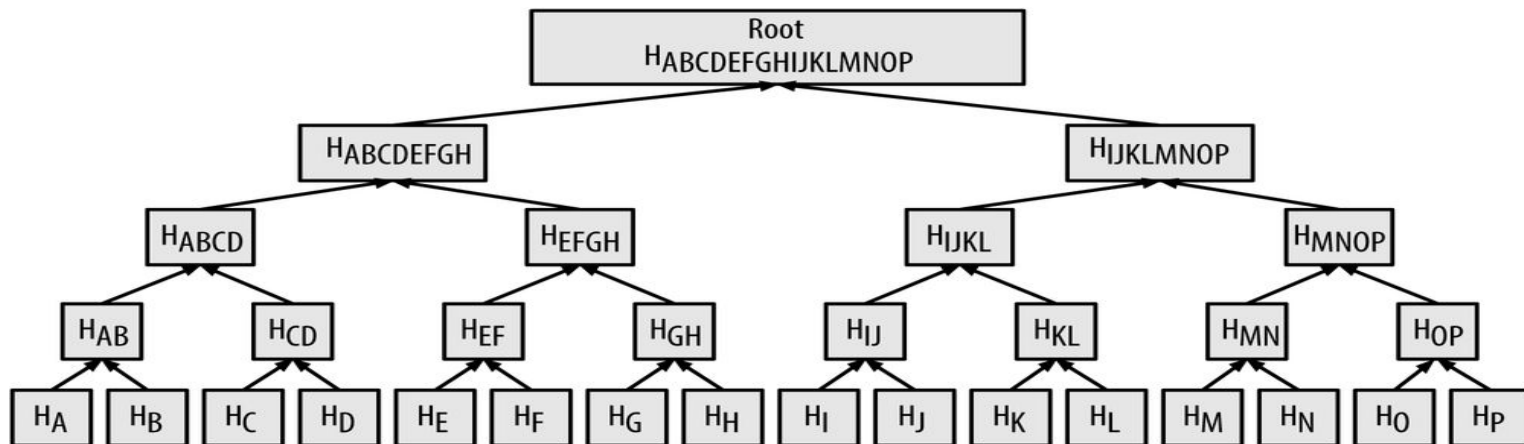
$$H_B = \text{SHA256}(\text{SHA256}(\text{TxB}))$$



$$H_{AB} = \text{SHA256}(\text{SHA256}(H_A + H_B))$$

# Introduction - Merkle Tree Structure

- A typically Merkle tree has the structure similar to the figure depicted below
  - Please note that the hash functions are double SHA-256

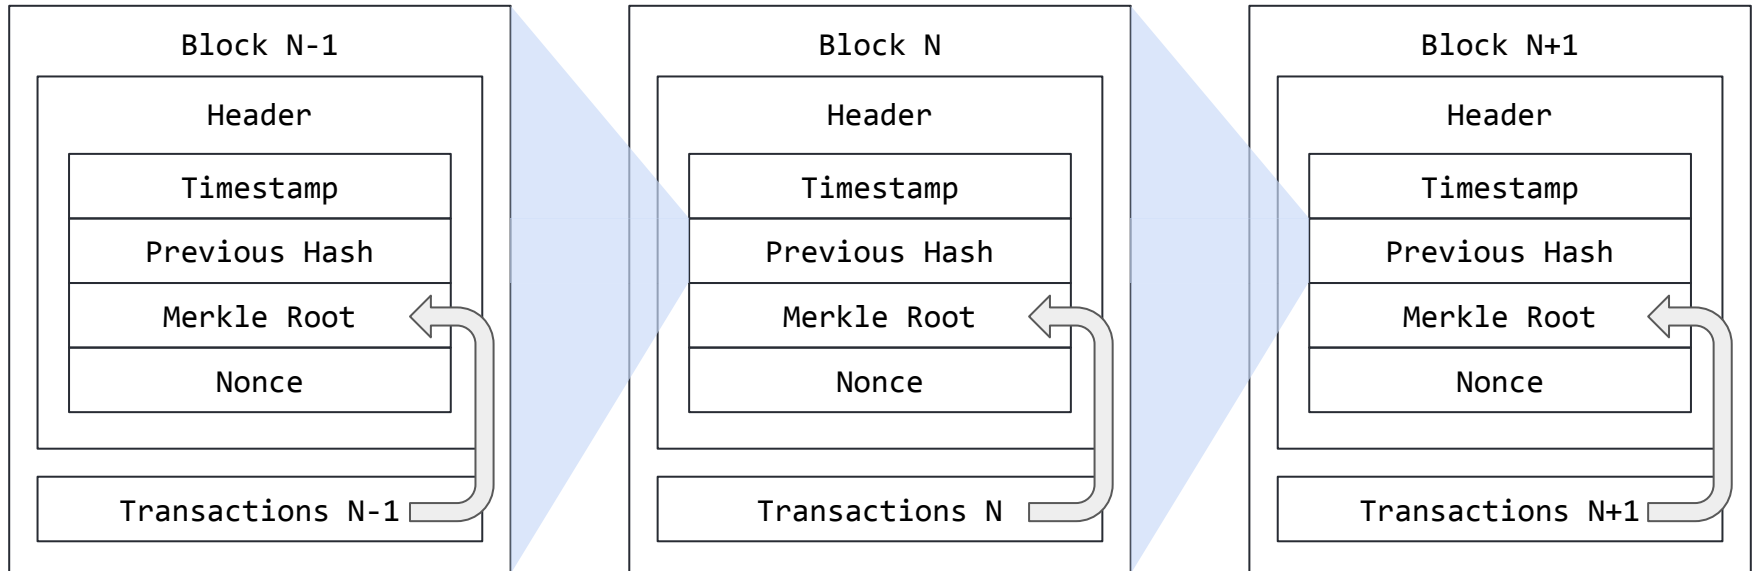


Credit: <https://www.oreilly.com/library/view/mastering-bitcoin/9781491902639/ch07.html>



# Introduction - Blockchain

- A blockchain is essentially a sequence of verified blocks
  - Each block contains the has value of its previous block's header
  - If something is modified, it will be immediately discovered



# Introduction - Verification of Transactions

- Who has to verify these transactions?
  - The nodes in the bitcoin network
- How do they verify the transactions?
  - The transactions from each node are broadcast to the entire network
  - Each node collects transactions by some algorithms defined by the bitcoin protocol.
  - The transactions are collected for block encryption
- How is the encrypted block added to the blockchain?
  - The newly encrypted block is broadcast to the entire network
  - The miners in the network verifies if the new block is correct and conforms to the bitcoin protocol (e.g., Whether block header's hash < Target Difficulty or not)

# Introduction - Target Difficulty

- The Target Difficulty is a value to determine how difficult it is to find out a proper hash value for the block header
- The Target Difficulty conforms to the following rules
  - It is a 256-bit number, encoded by a field called “nbits” (in 4 bytes)
  - It is determined by the bitcoin network
  - It has n leading zeros (currently, n equals 76).
- The Target Difficulty is adjusted such that a new block is computed every 10 minutes (on average)
  - The value of n is determined by the computing power of the entire bitcoin network
  - It's updated every 2,016 blocks (roughly every two weeks).
  - You may check current difficulty here: <https://btc.com/stats/diff>

# Introduction - Goal

- The primary goal of assignment 4 is to find out a proper hash value for a given block header
  - The hash value has to be less than the Target Difficulty
  - The time you spent on finding the hash value has to be accelerated
  - Use any techniques that you've learned in this class!

# Sequential Code

1. Read input
2. Calculate a merkle root from transactions
3. Decode find out the Target Difficulty
4. Find out a proper nonce, such that the hash value of the block header satisfies the requirement  
for nonce = 0X00000000 to 0Xffffffff  
    calculate a hash value  
    if the hash value < Target Difficulty  
        done

# Sequential Code - Header Extraction

- Step 1: Read input
- There are several fields in the block header that you have to extract.
- Please familiar yourself with the fields extracted from the header.

```
// **** read data ****
char version[9];
char prevhash[65];
char ntime[9];
char nbits[9];
int tx;
char *raw_merkle_branch;
char **merkle_branch;

getline(version, 9, fin);
getline(prevhash, 65, fin);
getline(ntime, 9, fin);
getline(nbits, 9, fin);
fscanf(fin, "%d\n", &tx);
printf("start hashing")

raw_merkle_branch = new char [tx * 65];
merkle_branch = new char *[tx];
for(int i=0;i<tx;++i) {
    merkle_branch[i] = raw_merkle_branch + i * 65;
    getline(merkle_branch[i], 65, fin);
    merkle_branch[i][64] = '\0';
}
```

# Sequential Code - Header Fields

- A representative header fields of a block is depicted below
  - Please have a comparison with pages 6, 10, and 16
  - These information has to be extracted first before we carrying out the calculation
- **case01.in**

4	Number of blocks
20000000	Version
000000000000000000003348540cbfc68b70825e7abcd5a83a48a5f87fa7f1aace	Previous block hash
5ac22f8b	Timestamp
17502ab7	Bits (packed difficulty)
2094	Number of transactions
c6574adb277efbfb972658ab78b1277707a967076cfc90d6af800cd8a915396d	Transactions
c01c33240ba97fb6db2b98fbaf7e4211fe3b59585372994bdac1b96b1c9be0d3	
abce7b2b30ff62b4998864db6a1ea77db8eb33d3f6abe3f981e0d2802c999042	
fd463aca59df1302c8f08e8070b56bd64ebe0cf35ba68cd5b39d208a2ff50053	
4321318516cb6630e3b3caa29bd49227168a69d170e1110c220b3c30d15f913c	
f1f452e09dad935a67499a24e388d61fd7f0f38c97a2a92c73b2fd7019a85578	
dd24a8e3f419006cd2f7cb3336b605fa6c2898accecac3d0a844c2ae8f5f53d9	

# Sequential Code - Merkle Root

- Step 2: Calculate a merkle root from transactions
  - The merle root is calculated by a function provided by the TA
  - You are encouraged to take a look of its implementation

```
// **** calculate merkle root ****  
  
unsigned char merkle_root[32];  
calc_merkle_root(merkle_root, tx, merkle_branch);  
  
printf("merkle root(little): ");  
print_hex(merkle_root, 32);  
printf("\n");  
  
printf("merkle root(big):      ");  
print_hex_inverse(merkle_root, 32);  
printf("\n");
```



# Sequential Code - Target Difficulty

- Step 3: Decode find out the Target Difficulty

- The is decoded and calculated from a number called bits
- The decode algorithm is already implemented by us (shown below)
- If you are interested in that algorithm, please take a look of this [website](#).

```
// ***** calculate target value *****  
// calculate target value from encoded difficulty  
// which is encoded on "nbits"  
unsigned int exp = block.nbits >> 24;  
unsigned int mant = block.nbits & 0xffffffff;  
unsigned char target_hex[32] = {};  
  
unsigned int shift = 8 * (exp - 3);  
unsigned int sb = shift / 8;  
unsigned int rb = shift % 8;
```

- Example:

- 2022/04/14 Bits: 0x1709f8d9
- Target =  $0x09f8d9 * 2^{(8 * (0x17 - 3))}$   
= 0x0000000000000000000000009d89700000000000000000000000000000000000

# Sequential Code - Nonce

- Step 4 : Find out a proper nonce, such that the hash value of the block header satisfies the requirement
  - Try if you can figure out a way to parallelize and accelerate the code

```
for(block.nonce=0x00000000; block.nonce<=0xffffffff;++block.nonce) {  
    //sha256d  
    double_sha256(&sha256_ctx, (unsigned char*)&block, sizeof(block));  
    if(block.nonce % 1000000 == 0) {  
        printf("hash %#10u (big): ", block.nonce);  
        print_hex_inverse(shash256_ctx.b, 32);  
        printf("\n");  
    }  
    if(little_endian_bit_comparison(shash256_ctx.b, target_hex, 32) < 0) { // shash256_ctx < target_hex  
        printf("Found Solution!!\n");  
        printf("hash %#10u (big): ", block.nonce);  
        print_hex_inverse(shash256_ctx.b, 32);  
        printf("\n\n");  
        break;  
    }  
}
```

# Sequential Code - SHA-256

- SHA-256 implementation for the sequential version can be found at:
  - sha256.h
  - sha256.cu
- If you are interested in the implantation details, please take a look at the following website
  - <https://en.wikipedia.org/wiki/SHA-2>

# Testcases

- We provide several testcases for you to test your program
- The execution time of the sequential code are given below:
  - case00.in : about 1 hr 15 min
  - case01.in : about 1 hr 15 min
  - case02.in : about 30 min
  - case03.in : about 11 min
- Please keep in mind that your goal is to parallelize the mining algorithm and accelerating it.
- It has to be faster than the sequential version.

# Grading

- **Correctness (40%)**
  - Please make sure to use GPU and verify your answer
- **Performance (30%)**
  - We will grade the performance of your assignments against the other students in the class
- **Report (30%)**
  - Please describe, in detail, how you parallelize your codes as well as your optimization methodology
- **Advanced CUDA skills**
  - You are welcome to use streaming, page-lock memory, asynchronous memory copy, or any other advanced skills
- **Others**
  - You will get credits if you are able to optimize the other parts of the source codes. Please justify your solutions and provide a detailed comparison in your report

# Grading - Report

- In your report, please include the following parts
  - Your implementation
  - The parallelization and optimization techniques you used in your solution
  - Experiments of various combinations of the number of blocks & threads (at least 8 combinations) and plot them with the figures
  - Describe the details if you use advanced CUDA skills
  - If you optimize the other parts of your source codes, please demonstrate your experimental results. We REQUIRE you to justify your solutions so that we can give you credits.

# Submission

- Please submit your hw4 to eeclass
  - hw4.cu
  - sha256.cu
  - sha256.h
  - report.pdf
  - Makefile (optional)
- Please do not package them, directly upload the files to eeclass
- Please make sure your Makefile works properly and your program can run before submitting your assignment 4

# Reminder

- Because we are doing hash, there may be a situation multiple nonce can satisfies the requirement  $\text{hash value} < \text{Target Difficulty}$
- We accept all nonce satisfies the requirement, so don't worry if your nonce isn't same as our provided solutions
- We ensure that all the testcases have more than one solution
- Your Makefile should build executable binary hw4



# Deadline

- The deadline of assignment 4 is 5/10 (Tue), 23:59
- Everyone is welcomed to ask questions on eeclass