

Lab3


GPU Cluster

Introduction to Parallel Computing
2022/03/22

Hardware

- **hades01**
 - CPU: Intel Xeon X5670 (2.93GHz) * 2
 - Memory: 96GB
 - No GPU
- **hades02 - 07**
 - CPU: ---
 - Memory: 64GB
 - GPU: GTX 1080*2

Architecture



You **cannot** use GPU on hades03-07 directly !

Software

- Linux 3.10
- Compilers
 - GCC 10.1.0
 - Clang 10.0.1
 - NVCC 11.0.194 (CUDA 11)
- Wordload Scheduler
 - Slurm 20.02.5

Available Resources

- Login node (200% CPU max)
- hades02
 - 2 GPUs, use by `export CUDA_VISIBLE_DEVICES` to env
- The remaining are compute nodes, available via job submission
 - Max nodes: 2
 - Max GPU: 2
 - Max wall time: 5 minutes
 - Max jobs run at any time: 2
- Priority
 - Favor short running jobs(based on wall time)
 - Favor less resource demanding jobs (based on number of GPUs)
 - Favor jobs which are queued for longer

SSH Credentials

- Hostname
 - `hades.cs.nthu.edu.tw`
- Username
 - Same as apollo
- Password
 - Same as apollo
- Check the slides of Lab1 if you forget how to login to the server
 - If you don't want to re-type your password when ssh into hades02
=> Generate the public/private key pair on hades, and ssh-copy-id into hades!

Compile A CUDA Program

- Compile a CUDA program
 - `nvcc [options] [exe] [target]`
 - e.g. `nvcc -o exe src.cu`
- Most of usages are same as GNU compiler
- Check `nvcc --help` for more options

hades02

- ssh to hades02
- export CUDA_VISIBLE_DEVICES=0 for using the first GPU
- export CUDA_VISIBLE_DEVICES=1 for using the second GPU
- export CUDA_VISIBLE_DEVICES=0,1 for using the two GPUs
- Run program on hades02 directly

```
engine210@hades02 ~-> export CUDA_VISIBLE_DEVICES=0,1
engine210@hades02 ~-> nvidia-smi
Mon Mar 21 21:12:00 2022

+-----+
| NVIDIA-SMI 450.57          Driver Version: 450.57          CUDA Version: 11.0          |
+-----+-----+
| GPU   Name           Persistence-M| Bus-Id        Disp.A | Volatile Uncorr. ECC |
| Fan  Temp  Perf    Pwr:Usage/Cap|      Memory-Usage | GPU-Util  Compute M. |
|                                           MIG M.         |
+-----+-----+
| 0     GeForce GTX 1080    On      | 00000000:4B:00:0 Off |           N/A       |
| 0%   38C    P8       7W / 200W    | 1MiB / 8119MiB      |           0%      Default |
|                                           N/A               |
+-----+-----+
| 1     GeForce GTX 1080    On      | 00000000:4D:00:0 Off |           N/A       |
| 0%   45C    P8      14W / 200W    | 1MiB / 8117MiB      |           0%      Default |
|                                           N/A               |
+-----+-----+

+-----+
| Processes: |
| GPU   GI   CI        PID   Type   Process name                      GPU Memory |
| ID     ID                                   |              Usage |
+-----+-----+
| No running processes found |
+-----+
```


Slurm Scheduler (For hades[03-07])

- Most of options are same as the apollo scheduler

-p <PARTITION>	<PARTITION> should be ipc22
-N <NODES>	<NODES> should be the number of nodes to run the job
-n <PROCESSES>	<PROCESSES> are the number of processes to launch
-c <CORES>	Each process can use up to <CORES> CPU cores
-t <TIME>	Time limit with format “mm:ss”
-j <NAME>	The name of job, will be displayed on queue
-w <NODE_LIST>	The specific <NODE_LIST> you want to run
--gres=gpu:<GPUS>	<GPUS> is the numbers of GPUs to run the job

Example

```
[ipc22sxx@hades01 ~]$ srun -p ipc22 --gres=gpu:1 -N 1 -n 1 -w hades04 ./deviceQuery
```

Demo

- [CUDA-MEMCHECK](#)

```
cuda-memcheck /home/ipc22/share/lab3/memcheck_demo
```

- [CUDA-GDB](#)

Not available on hades, but you can try it on your own computer.

Practice

- Google Form: <https://forms.gle/DHteQJ44vQh943mW8>
- Please finish the form before 3/29 23:59

Note for VSCode Terminal users. If command not found, try:

- `bash -l`
- MPI on apollo (mpicc, mpirun, etc.):
`source /opt/intel/mpi/intel64/bin/mpivars.sh`
- CUDA on hades (nvcc, etc.):
`source /etc/profile.d/cuda.sh`